

NC State University
Department of Electrical and Computer Engineering

ECE 463/521: Fall 2014

Project #2: Branch Prediction (Version 1.1)

Due: 11:59pm 11/13/14

1. Ground rules:

1. All students must work alone. The project scope is reduced (but still substantial) for ECE 463 students, as detailed in this specification.
2. Sharing of code between students is considered cheating and will receive appropriate action in accordance with University policy. The TAs will scan source code (from current and past semesters) through various tools available to us for detecting cheating. Source code that is flagged by these tools will be dealt with severely.
3. A Wolfware message board will be provided for posting questions, discussing and debating issues, and making clarifications. It is an essential aspect of the project and communal learning. Students must not abuse the message board, whether inadvertently or intentionally. Above all, never post actual code on the message board (unless permitted by the TAs/instructor). When in doubt about posting something of a potentially sensitive nature, email the TAs and instructor to clear the doubt.
4. You *must* do all your work in the C/C++ or Java languages. Exceptions must be approved. The C language is fine to use, as that is what many students are trained in. Basic C++ extensions to the C language (*e.g.*, classes instead of structs) are encouraged (but by no means required) because it enables more straightforward code-reuse.
5. Use of the Eos Linux environment is *required*. This is the platform where the TAs will compile and test your simulator. (WARNING: If you develop your simulator on another platform, get it working on that platform, and then try to port it over to Eos Linux at the last minute, you may encounter major problems. Porting is not as quick and easy as you think unless you are an excellent programmer. What's worse, malicious bugs can be hidden until you port the code to a different platform, which is an unpleasant surprise close to the deadline.)

2. Project Description

In this project, you will construct a branch predictor simulator and use it to design branch predictors well suited to the SPECint95 benchmarks.

3. Simulator Specification

3.1. ECE 463 and ECE 521 students: Branch predictors

Model a *gshare* branch predictor with parameters $\{m, n\}$, where:

- m is the number of low-order PC bits used to form the prediction table index. **Note:** discard the lowest two bits of the PC, since these are always zero, i.e., use bits $m+1$ through 2 of the PC.
- n is the number of bits in the global branch history register. **Note:** $n \leq m$. **Note:** n may equal zero, in which case we have the simple *bimodal* branch predictor.

3.1.1. $n=0$: bimodal branch predictor

When $n=0$, the *gshare* predictor reduces to a simple *bimodal* predictor. In this case, the index is based on only the branch's PC, as shown in Fig. 1 below.

Entry in the prediction table:

An entry in the prediction table contains a single 2-bit counter. All entries in the prediction table should be initialized to **2** ("weakly taken") when the simulation begins.

Regarding branch interference:

Different branches may index the same entry in the prediction table. This is called "interference". Interference is not explicitly detected or avoided: it just happens. (There is no tag array, no tag checking, and no "miss" signal for the prediction table!)

Steps:

When you get a branch from the trace file, there are three steps:

- (1) Determine the branch's **index** into the prediction table.
- (2) Make a prediction. Use **index** to get the branch's counter from the prediction table. If the counter value is greater than or equal to **2**, then the branch is predicted *taken*, else it is predicted *not-taken*.
- (3) Update the branch predictor based on the branch's actual outcome. The branch's counter in the prediction table is incremented if the branch was taken, decremented if the branch was not-taken. The counter *saturates* at the extremes (**0** and **3**), however.




Figure 1. Bimodal branch predictor.

3.1.2. $n > 0$: gshare branch predictor

When $n > 0$, there is an n -bit global branch history register. In this case, the index is based on both the branch's PC and the global branch history register, as shown in Fig. 2 below. The global branch history register is initialized to all zeroes (00...0) at the beginning of the simulation.

Steps:

When you get a branch from the trace file, there are three steps:

- (1) Determine the branch's **index** into the prediction table. Fig. 2 shows how to generate the index: the current n -bit global branch history register is XORed with the uppermost n bits of the index(PC) bits.
- (2) Make a prediction. Use **index** to get the branch's counter from the prediction table. If the counter value is greater than or equal to **2**, then the branch is predicted *taken*, else it is predicted *not-taken*.
- (3) Update the branch predictor based on the branch's actual outcome. The branch's counter in the prediction table is incremented if the branch was taken, decremented if the branch was not-taken. The counter *saturates* at the extremes (**0** and **3**), however.
- (4) Update the global branch history register. Shift the register right by 1 bit position, and place the branch's actual outcome into the most-significant bit position of the register.

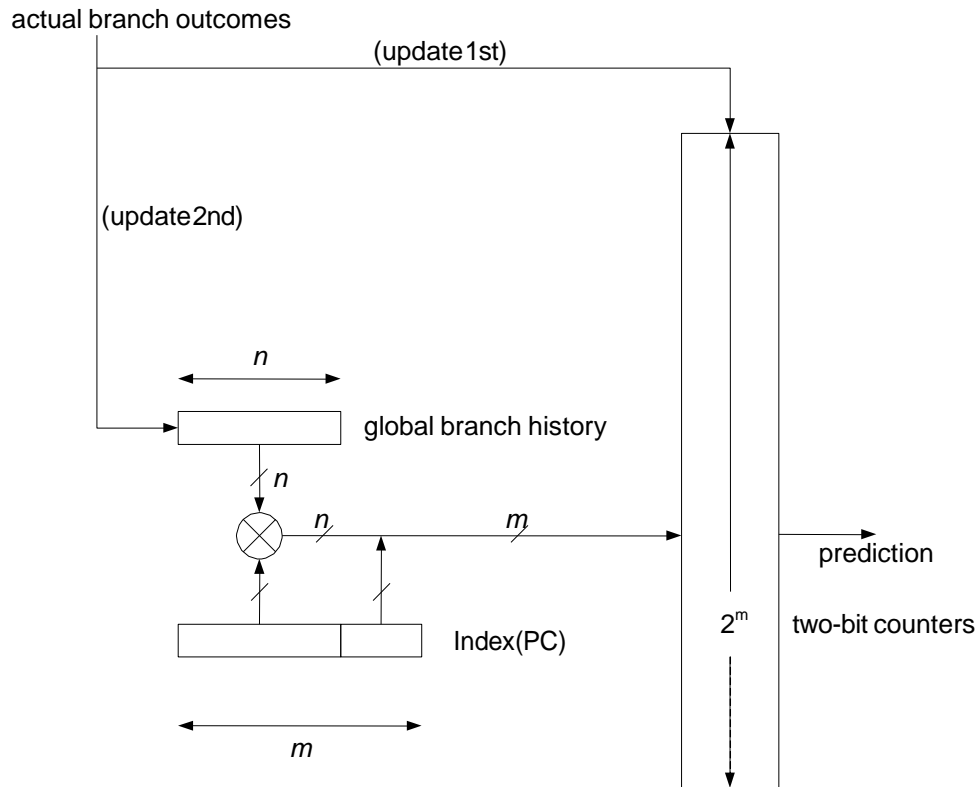


Figure 2. Gshare branch predictor.

3.2. ECE 521 students only: Hybrid branch predictor

Model a **hybrid predictor** that selects between the *bimodal* and the *gshare* predictors, using a chooser table of 2^k 2-bit counters. All counters in the chooser table are initialized to **1** at the beginning of the simulation.

Steps:

When you get a branch from the trace file, there are six top-level steps:

- (1) Obtain two predictions, one from the *gshare* predictor (follow steps 1 and 2 in Section 3.1.2) and one from the *bimodal* predictor (follow steps 1 and 2 in Section 3.1.1).
- (2) Determine the branch's **index** into the chooser table. The index for the chooser table is bit $k+1$ to bit **2** of the branch PC (*i.e.*, as before, discard the lowest two bits of the PC).
- (3) Make an overall prediction. Use **index** to get the branch's chooser counter from the chooser table. If the chooser counter value is greater than or equal to **2**, then use the prediction that was obtained from the *gshare* predictor, otherwise use the prediction that was obtained from the *bimodal* predictor.
- (4) Update the selected branch predictor based on the branch's actual outcome. Only the branch predictor that was selected in step 3, above, is updated (if *gshare* was selected, follow step 3 in Section 3.1.2, otherwise follow step 3 in Section 3.1.1).
- (5) Note that the *gshare*'s global branch history register must always be updated, even if *bimodal* was selected (follow step 4 in Section 3.1.2).
- (6) Update the branch's chooser counter using the following rule:

	Results from predictors:		
	<i>both incorrect or both correct</i>	<i>gshare correct, bimodal incorrect</i>	<i>bimodal correct, gshare incorrect</i>
Chooser counter update policy:	no change	increment (but saturates at 3)	decrement (but saturates at 0)

3.3. Bonus part: Branch Target Buffer (optional)

Model a *set-associative* **branch target buffer (BTB)** that stores information about previously seen branches, ignoring the lower two bits of the PC as they are zeros for an ISA with the instruction size of 4 bytes. (as shown in Fig. 3)

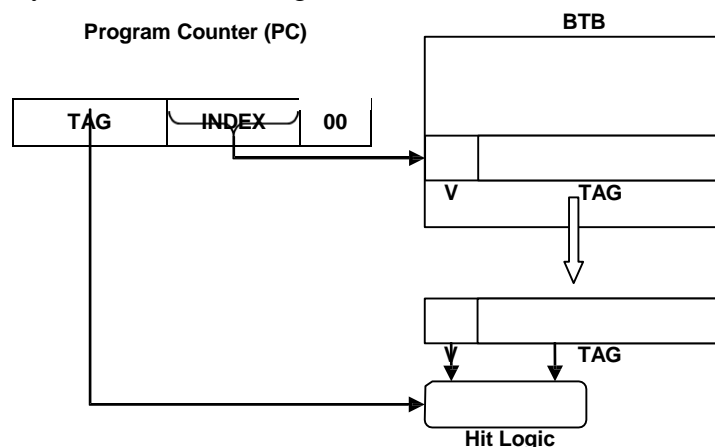


Figure 3. Branch target buffer

When you get a branch from the trace file, the BTB is first checked for a hit. There are two possible scenarios:

1. **Scenario #1: Branch misses in BTB.** In this case, the branch is implicitly predicted as not taken. The following two step is performed:
 - i) Make space for the new PC. If there is at least one invalid block in the set, then there is already space for the PC and no further action is required (go to step ii). On the other hand, if all blocks in the set are valid, then a victim block V must be singled out for eviction, according to the **LRU** replacement policy. No writeback is needed upon eviction.
 - ii) Add the PC into branch target buffer.
 If the branch turned out to be actually taken, it will be considered as a misprediction.
2. **Scenario #2: Branch hits in BTB.** In this case, nothing happens with respect to branch target buffer. Predictions will be made from branch predictors. Follow the steps in Section 3.1 and 3.2 to complete the branch prediction. Note that predictions from branch predictors are only made when the branch hits in BTB.

HINT:

You can reuse your code from Project 1-A (L1 cache) to model the BTB. To simplify the simulation, you can set the block size to be 4 as we are ignoring the last two bits.

4. Inputs to Simulator

The simulator reads a trace file in the following format:

```
<hex branch PC> t|n
<hex branch PC> t|n
...
```

Where:

- o <hex branch PC> is the address of the branch instruction in memory. This field is used to index into predictors.
- o "t" indicates the branch is actually taken (Note! *Not* that it is predicted taken!). Similarly, "n" indicates the branch is actually not-taken.

Example:

```
00a3b5fc t
00a3b604 t
00a3b60c n
...
```

5. Outputs from Simulator

The simulator outputs the following measurements after completion of the run:

Without BTB:

- a. number of accesses to the predictor (*i.e.*, number of branches)
- b. number of branch mispredictions (predicted *taken* when *not-taken*, or predicted *not-taken* when *taken*)
- c. branch misprediction rate (# mispredictions/# branches)

With BTB:

- a. size of BTB
- b. number of branches
- c. number of predictions from Branch Predictor (number of BTB hits)
- d. number of mispredictions from Branch Predictor (predicted *taken* when *not-taken*, or predicted *not-taken* when *taken*)
- e. number of branches miss in BTB and taken (number of mispredictions due to BTB miss)
- f. total number of mispredictions (item d + item e)
- g. branch misprediction rate (item f / item b)

6. Validation and Other Requirements

6.1. Validation requirements

Sample simulation outputs will be provided on the website. These are called “validation runs”. You must run your simulator and debug it until it matches the validation runs.

Each validation run includes:

1. The branch predictor configuration.
2. The final contents of the branch predictor.
3. All measurements described in Section 5.

Your simulator must print outputs to the console (i.e., to the screen). (Also see Section 6.2 about this requirement.)

Your output must match both numerically and in terms of formatting, because the TAs will literally “diff” your output with the correct output. You must confirm correctness of your simulator by following these two steps for each validation run:

- 1) Redirect the console output of your simulator to a temporary file. This can be achieved by placing `> your_output_file` after the simulator command.
- 2) Test whether or not your outputs match properly, by running this unix command:

```
diff -iw <your_output_file> <posted_output_file>
```

The `-iw` flags tell “diff” to treat upper-case and lower-case as equivalent and to ignore the amount of whitespace between words. Therefore, you do not need to worry about the exact number of spaces or tabs as long as there is some whitespace where the validation runs have whitespace.

6.2. Compiling and running simulator

You will hand in source code and the TAs will compile and run your simulator. As such, you must meet the following strict requirements. Failure to meet these requirements will result in point deductions (see section “Grading”).

1. You must be able to compile and run your simulator on Eos Linux machines. This is required so that the TAs can compile and run your simulator. If you are logging into an Eos machine remotely and do not know whether or not it is Linux (as opposed to SunOS), use the “uname” command to determine the operating system.
2. Along with your source code, you must provide a **Makefile** that automatically compiles the simulator. This Makefile must create a simulator named “sim”. The TAs should be able to type only “make” and the simulator will successfully compile. The TAs should be able to type only “make clean” to automatically remove object (.o) files and the simulator executable. An example Makefile will be posted on the web page, which you can copy and modify for your needs.

3. Your simulator must accept command-line arguments as follows:

To simulate a bimodal predictor: **sim bimodal <M2> <BTB size> <BTB assoc> <tracefile>**, where M2 is the number of PC bits used to index the bimodal table; BTB assoc is the associativity of the BTB, BTB size and BTB assoc are 0 if no BTB is modeled.

To simulate a gshare predictor: **sim gshare <M1> <N> <BTB size> <BTB assoc> <tracefile>**, where M1 and N are the number of PC bits and global branch history register bits used to index the gshare table, respectively; BTB assoc is the associativity of the BTB, BTB size and BTB assoc are 0 if no BTB is modeled.

To simulate a hybrid predictor: **sim hybrid <K> <M1> <N> <M2> <BTB size> <BTB assoc> <tracefile>**, where K is the number of PC bits used to index the chooser table, M1 and N are the number of PC bits and global branch history register bits used to index the gshare table (respectively), and M2 is the number of PC bits used to index the bimodal table. BTB assoc is the associativity of the BTB, BTB size and BTB assoc are 0 if no BTB is modeled.

(<tracefile> is the filename of the input trace.)

4. Your simulator must print outputs to the console (i.e., to the screen). This way, when a TA runs your simulator, he/she can simply redirect the output of your simulator to a filename of his/her choosing for validating the results.

6.3. Keeping backups

It is good practice to frequently make backups of all your project files, including source code, your report, etc. You can backup files to another hard drive (Eos account vs. home PC vs. laptop ... keep consistent copies in multiple places) or removable media (Flash drive, etc.).

6.4. Run time of simulator

Correctness of your simulator is of paramount importance. That said, making your simulator *efficient* is also important for a couple of reasons.

First, the TAs need to test every student's simulator. Therefore, we are placing the constraint that your simulator must finish a single run in 2 minutes or less. If your simulator takes longer than 2 minutes to finish a single run, please see the TAs as they may be able to help you speed up your simulator.

Second, you will be running many experiments: many branch predictor configurations and multiple traces. Therefore, you will benefit from implementing a simulator that is reasonably fast.

One simple thing you can do to make your simulator run faster is to compile it with a high optimization level. The example Makefile posted on the web page includes the `-O3` optimization flag.

Note that, when you are debugging your simulator in a debugger (such as gdb), it is recommended that you compile without `-O3` and with `-g`. Optimization includes register allocation. Often, register-allocated variables are not displayed properly in debuggers, which is why you want to disable optimization when using a debugger. The `-g` flag tells the compiler to include symbols (variable names, etc.) in the compiled binary. The debugger needs this information to recognize variable names, function names, line numbers in the source code, etc. When you are done debugging, recompile with `-O3` and without `-g`, to get the most efficient simulator again.

6.5. VCL Information

In addition to using *grendel cluster*, and other eos linux machines, VCL allows you to reserve linux machines. The following website describes how to use VCL:

<http://vcl.ncsu.edu/>

From the website:

- 1) select “Reservation System > Make a Reservation”,
- 2) select “NCSU WRAP” and “Proceed to Login”,
- 3) login,
- 4) for the environment, select “Linux Lab Machine (Realm RHEnterprise 4)” or “Linux Lab Machine (Realm RHEnterprise 5)”,
- 5) you may choose to reserve a shell for up to 4 hours beginning now or later (with the possibility to extend reservation before it expires),
- 6) click “Create Reservation” button,
- 7) wait until the screen updates with a reservation and click “Connect!”,
- 8) you will be given an internet address that you can ssh to, using putty or other ssh clients, the same way that you remotely and securely login to other linux machines.

7. Tasks, Grading Breakdown

PART 1: BIMODAL PREDICTOR

(a) [ECE463: 25 points] or [ECE521: 20 points] Match the four validation runs “val_bimodal_1.txt”, “val_bimodal_2.txt”, “val_bimodal_3.txt”, and “val_bimodal_4.txt”, posted on the website for the BIMODAL PREDICTOR. Each validation run is worth $\frac{1}{4}$ of the points for this part (6 or 5 points each). You must match all validation runs to get credit for the experiments with the bimodal predictor, however.

(b) [ECE463: 25 points] or [ECE521: 20 points] Simulate BIMODAL PREDICTOR configurations for $7 \leq m \leq 12$. Use the traces in the trace directory.

[15 or 10 points] Graphs: Produce one graph for each benchmark. Graph title: “<benchmark>, bimodal”. Y-axis: branch misprediction rate. X-axis: m . Per graph, there should be only one curve consisting of 6 datapoints (connect the datapoints with a line).

[5 points] Analysis: Draw conclusions and discuss trends. Discuss similarities/differences among benchmarks.

[5 points] Design: **For each trace**, choose a bimodal predictor design that minimizes misprediction rate AND minimizes predictor cost in bits. You have a maximum budget of **16 kilobytes** of storage. When “minimizing” misprediction rate, look for diminishing returns, *i.e.*, the point where misprediction rate starts leveling off and additional hardware provides only minor improvement. Bottom line: use good sense to provide high performance and reasonable cost. Justify your designs with supporting data and explanations.

PART 2: GSHARE PREDICTOR

(a) [ECE463: 25 points] or [ECE521: 20 points] Match the four validation runs “val_gshare_1.txt”, “val_gshare_2.txt”, “val_gshare_3.txt”, and “val_gshare_4.txt”, posted on the website for the GSHARE PREDICTOR. Each validation run is worth $\frac{1}{4}$ of the points for this part (6 or 5 points each). You must match all validation runs to get credit for the experiments with the gshare predictor, however.

(b) [ECE463: 25 points] or [ECE521: 20 points] Simulate GSHARE PREDICTOR configurations for $7 \leq m \leq 12$, $2 \leq n \leq m$, n is even. Use the traces in the trace directory.

[15 or 10 points] Graphs: Produce one graph for each benchmark. Graph title: “<benchmark>, gshare”. Y-axis: branch misprediction rate. X-axis: m . Per graph, there should be a total of 27 datapoints plotted as a family of 6 curves. Datapoints having the same value of n are connected with a line, *i.e.*, one curve for each value of n . Note that not all curves have the same number of datapoints.

[5 points] Analysis: Draw conclusions and discuss trends. Discuss similarities/differences among benchmarks.

[5 points] Design: **For each trace**, choose a gshare predictor design that minimizes misprediction rate AND minimizes predictor cost in bits. You have a maximum budget of **16 kilobytes** of storage. When “minimizing” misprediction rate, look for diminishing returns, *i.e.*,

the point where misprediction rate starts leveling off and additional hardware provides only minor improvement. Bottom line: use good sense to provide high performance and reasonable cost. Justify your designs with supporting data and explanations.

PART 3: HYBRID PREDICTOR (ECE521 students only)

[ECE521: 20 points] Match the two validation runs “val_hybrid_1.txt” and “val_hybrid_2.txt” posted on the website for the HYBRID PREDICTOR. The TAs will also check that your simulator matches two mystery validation runs. Each validation run is worth ¼ of the points for this part (5 points each).

PART 4: BRANCH TARGET BUFFER (bonus part)

[Bonus: 7/14 points for ECE521/ECE463] Match the two validation runs “val_BTBTB_1.txt” and “val_BTBTB_2.txt” posted on the website for the BTB. Each validation run is worth 3.5 or 7 points for this part for ECE521 and ECE463 students, respectively.

8. What to Submit via Wolfware

You must hand in a single zip file called **project2.zip** that is no more than 1MB in size. Please respect the size limit on behalf of fellow students, as the Wolfware submission space is limited and exceeding your quota will cause problems come submission time. (Notify the TAs beforehand if you have special space requirements. However, a zip file of 1MB should be sufficient.)

Below is an example showing how to create **project2.zip** from an Eos Linux machine. Suppose you have a bunch of source code files (*.cc, *.h), the Makefile, and your project report (report.doc).

```
zip project2 *.cc *.h Makefile report.doc
```

project2.zip must contain the following (any deviation from the following requirements may delay grading your project and may result in point deductions, late penalties, etc.):

1. **Project report.** This must be a single MS Word document named **report.doc** OR a single PDF document named **report.pdf**. The report must include the following:
 - A cover page with the project title, the Honor Pledge, and your full name as electronic signature of the Honor Pledge. A sample cover page will be posted on the project website.
 - See Section 7 for the required content of the report.
2. **Source code.** You must include the commented source code for the simulator program itself. You may use any number of .cc/.h files, .c/.h files, etc.

3. **Makefile.** See Section 6.2, item #2, for strict requirements. If you fail to meet these requirements, it may delay grading your project and may result in point deductions.

9. Penalties

Various deductions (out of 100 points):

-1 point for each hour late, according to Wolfware timestamp. **TIP:** Submit whatever you have completed by the deadline just to be safe. If, after the deadline, you want to continue working on the project to get more features working and/or finish experiments, you may submit a second version of your project late. If you choose to do this, the TA will grade both the on-time version and the late version (*late penalty applied to the late version*), and take the maximum score of the two. Hopefully you will complete everything by the deadline, however.

Up to -10 points for not complying with specific procedures. Follow all procedures very carefully to avoid penalties. Complete the **SUBMISSION CHECKLIST** that is posted on the project website, to make sure you have met all requirements.

Cheating: Source code that is flagged by tools available to us will be dealt with according to University Policy. This includes a 0 for the project and other disciplinary actions.